

lab2

Lillian Clark

9/6/2021

(Spent most of lab trying to install `rstanarm`)

```
bike <- read_csv("210830_bikecrash.csv")
```

```
##
## -- Column specification -----
## cols(
##   county = col_character(),
##   pop = col_double(),
##   med_hh_income = col_double(),
##   traffic_vol = col_double(),
##   pct_rural = col_double(),
##   crashes = col_double()
## )
```

```
bike_new <- bike %>%
  mutate(crashes_pc = crashes / (pop / 100000),
         high_crash_cty = case_when(
           crashes_pc >= 60 ~ 1,
           TRUE ~ 0),
         high_crash_cty = as.factor(high_crash_cty))
```


```
summary(bike_new$high_crash_cty)
```

```
##  0  1
## 63 37
```

```
summary(bike_new$pct_rural)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1.00  42.50   62.50   61.24  85.00  100.00
```

```
ggplot(bike_new, aes(x = high_crash_cty, y = pct_rural)) +
  geom_boxplot()
```



```
m0 <- glm(high_crash_cty ~ I(pop/1000000) + pct_rural, data = bike_new, family = "binomial")
round(summary(m0)$coef, 6)
```

```
##           Estimate Std. Error  z value Pr(>|z|)
## (Intercept)  0.994662   0.846443  1.175108 0.239951
## I(pop/1e+06)  1.162455   2.264275  0.513389 0.607679
## pct_rural    -0.028063   0.011431 -2.455100 0.014085
```

We use the estimate and standard error from a frequentist GLM as our prior for the intercept in our Bayesian regression.

```
m1 <- stan_glm(high_crash_cty ~ I(pop/1000000) + pct_rural,
               data = bike_new,
               family = binomial(link = "logit"),
               prior_intercept = normal(.995, .846),
               prior = normal(0, 100, autoscale = T),
               chains = 2, iter = 10000, seed = 9763,
               prior_PD = F)
```

```
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 0.000132 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 1.32 seconds.
## Chain 1: Adjust your expectations accordingly!
```

```

## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 10000 [ 0%] (Warmup)
## Chain 1: Iteration: 1000 / 10000 [ 10%] (Warmup)
## Chain 1: Iteration: 2000 / 10000 [ 20%] (Warmup)
## Chain 1: Iteration: 3000 / 10000 [ 30%] (Warmup)
## Chain 1: Iteration: 4000 / 10000 [ 40%] (Warmup)
## Chain 1: Iteration: 5000 / 10000 [ 50%] (Warmup)
## Chain 1: Iteration: 5001 / 10000 [ 50%] (Sampling)
## Chain 1: Iteration: 6000 / 10000 [ 60%] (Sampling)
## Chain 1: Iteration: 7000 / 10000 [ 70%] (Sampling)
## Chain 1: Iteration: 8000 / 10000 [ 80%] (Sampling)
## Chain 1: Iteration: 9000 / 10000 [ 90%] (Sampling)
## Chain 1: Iteration: 10000 / 10000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.399126 seconds (Warm-up)
## Chain 1:                0.491079 seconds (Sampling)
## Chain 1:                0.890205 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'bernoulli' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 2.3e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.23 seconds.
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 10000 [ 0%] (Warmup)
## Chain 2: Iteration: 1000 / 10000 [ 10%] (Warmup)
## Chain 2: Iteration: 2000 / 10000 [ 20%] (Warmup)
## Chain 2: Iteration: 3000 / 10000 [ 30%] (Warmup)
## Chain 2: Iteration: 4000 / 10000 [ 40%] (Warmup)
## Chain 2: Iteration: 5000 / 10000 [ 50%] (Warmup)
## Chain 2: Iteration: 5001 / 10000 [ 50%] (Sampling)
## Chain 2: Iteration: 6000 / 10000 [ 60%] (Sampling)
## Chain 2: Iteration: 7000 / 10000 [ 70%] (Sampling)
## Chain 2: Iteration: 8000 / 10000 [ 80%] (Sampling)
## Chain 2: Iteration: 9000 / 10000 [ 90%] (Sampling)
## Chain 2: Iteration: 10000 / 10000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.409604 seconds (Warm-up)
## Chain 2:                0.37632 seconds (Sampling)
## Chain 2:                0.785924 seconds (Total)
## Chain 2:

```

```
prior_summary(m1)
```


```

## Priors for model 'm1'
## -----
## Intercept (after predictors centered)
## ~ normal(location = 0.99, scale = 0.85)
##
## Coefficients
## Specified prior:

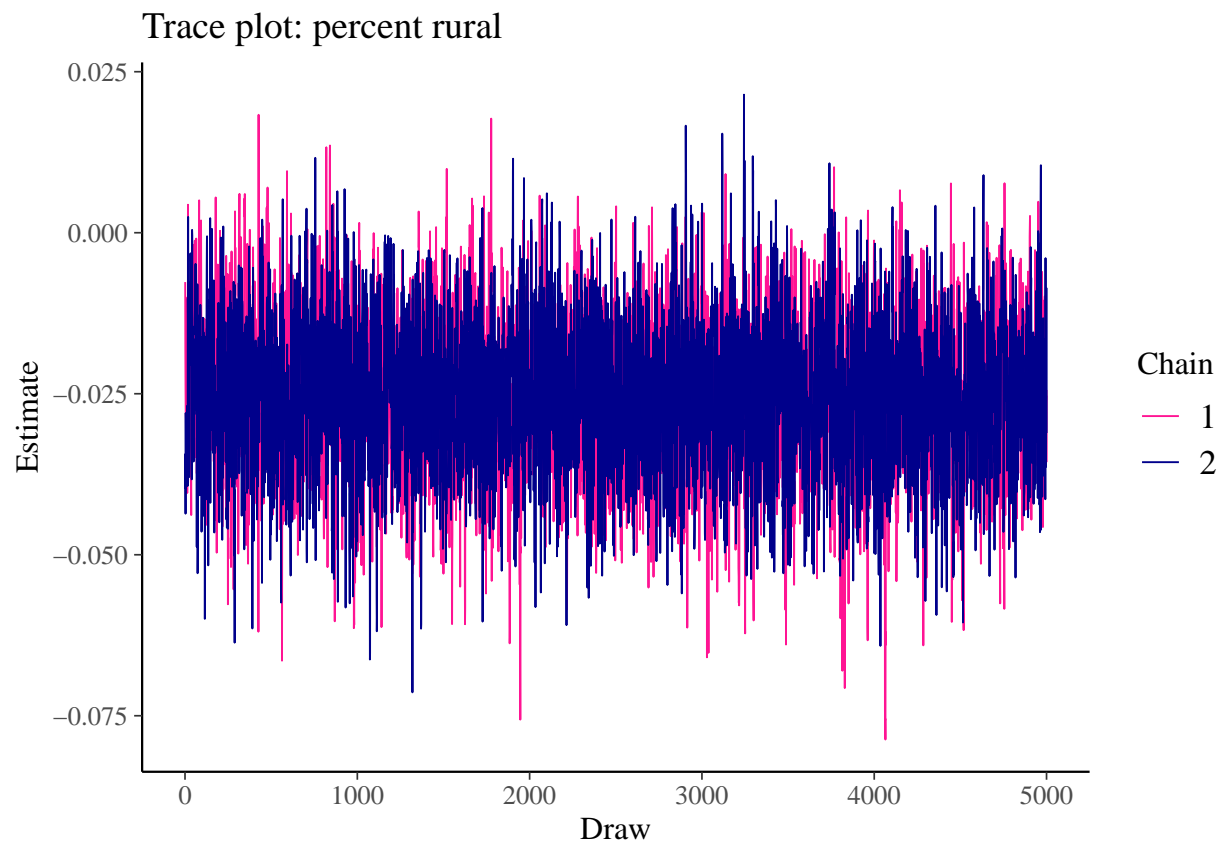
```

```
## ~ normal(location = [0,0], scale = [100,100])
## Adjusted prior:
## ~ normal(location = [0,0], scale = [597.13, 3.55])
## -----
## See help('prior_summary.stanreg') for more details
```

```
color_scheme_set(c("darkblue", "darkred", "darkgray",
                  "deepskyblue", "deeppink", "darkgreen"))
plot(m1, "trace")
```




```
plot(m1, "trace", pars = "pct_rural") +
  labs(title = "Trace plot: percent rural",
       y = "Estimate", x = "Draw")
```

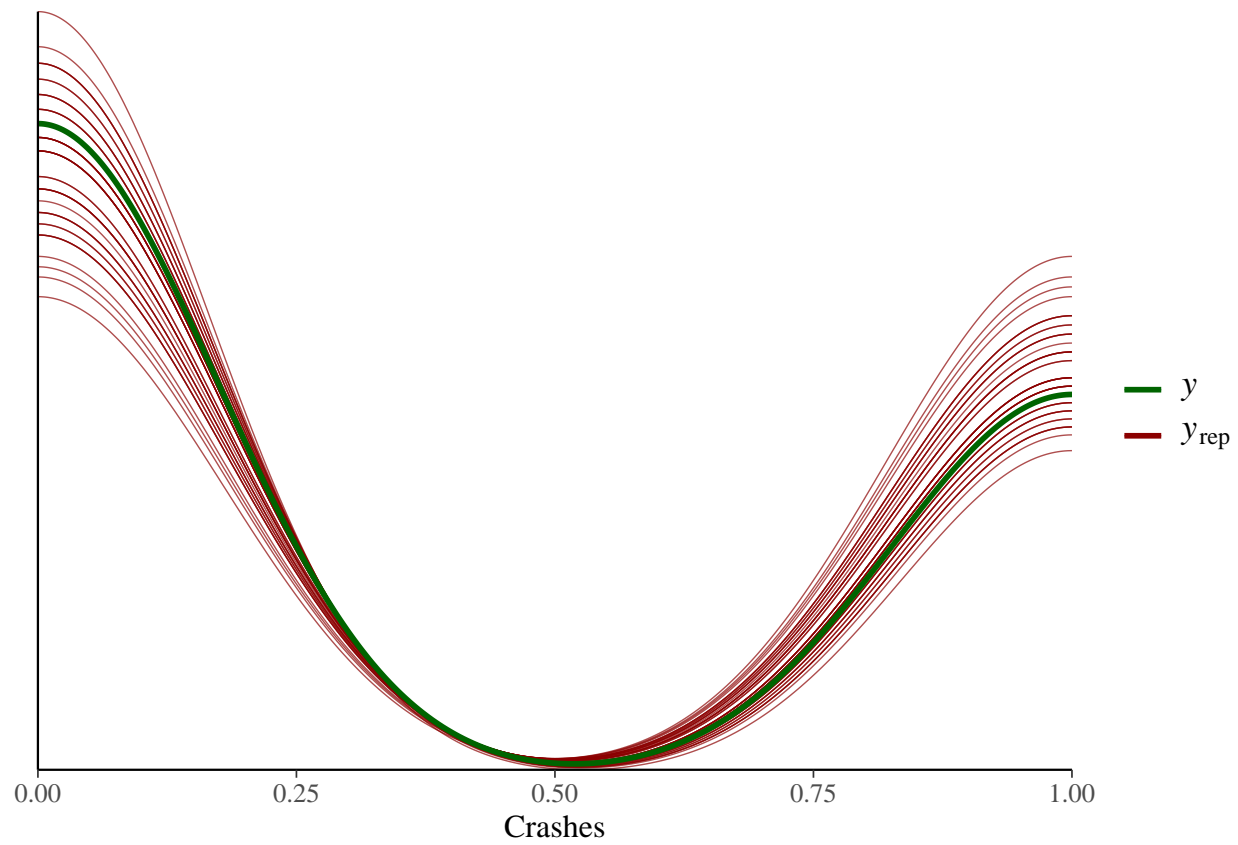


```
pp_check(m1, plotfun = "hist", nreps = 5) +  
  xlab("Crashes")
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
pp_check(m1) +
  xlab("Crashes")
```




```
plot(m1, "acf_bar") + # compare to "acf"  
  labs(title = "ACF plots")
```

ACF plots



```
plot(m1, "dens_overlay")
```

```
round(as.data.frame(summary(m1)), 2)
```

	mean	mcse	sd	10%	50%	90%	n_eff	Rhat
## (Intercept)	0.82	0.01	0.91	-0.36	0.83	1.95	4022	1
## I(pop/1e+06)	2.56	0.05	2.84	-0.73	2.22	6.38	3652	1
## pct_rural	-0.03	0.00	0.01	-0.04	-0.03	-0.01	4185	1
## mean_PPD	0.39	0.00	0.06	0.31	0.39	0.47	7221	1
## log-posterior	-63.47	0.02	1.31	-65.15	-63.12	-62.20	3372	1

```
bike_new %>%
  filter(county == "Durham")
```

```
## # A tibble: 1 x 8
##   county    pop med_hh_income traffic_vol pct_rural crashes crashes_pc
##   <chr>   <dbl>       <dbl>       <dbl>   <dbl>   <dbl>      <dbl>
## 1 Durham 316739      59.3         395       6      340      107.
## # ... with 1 more variable: high_crash_cty <fct>
```

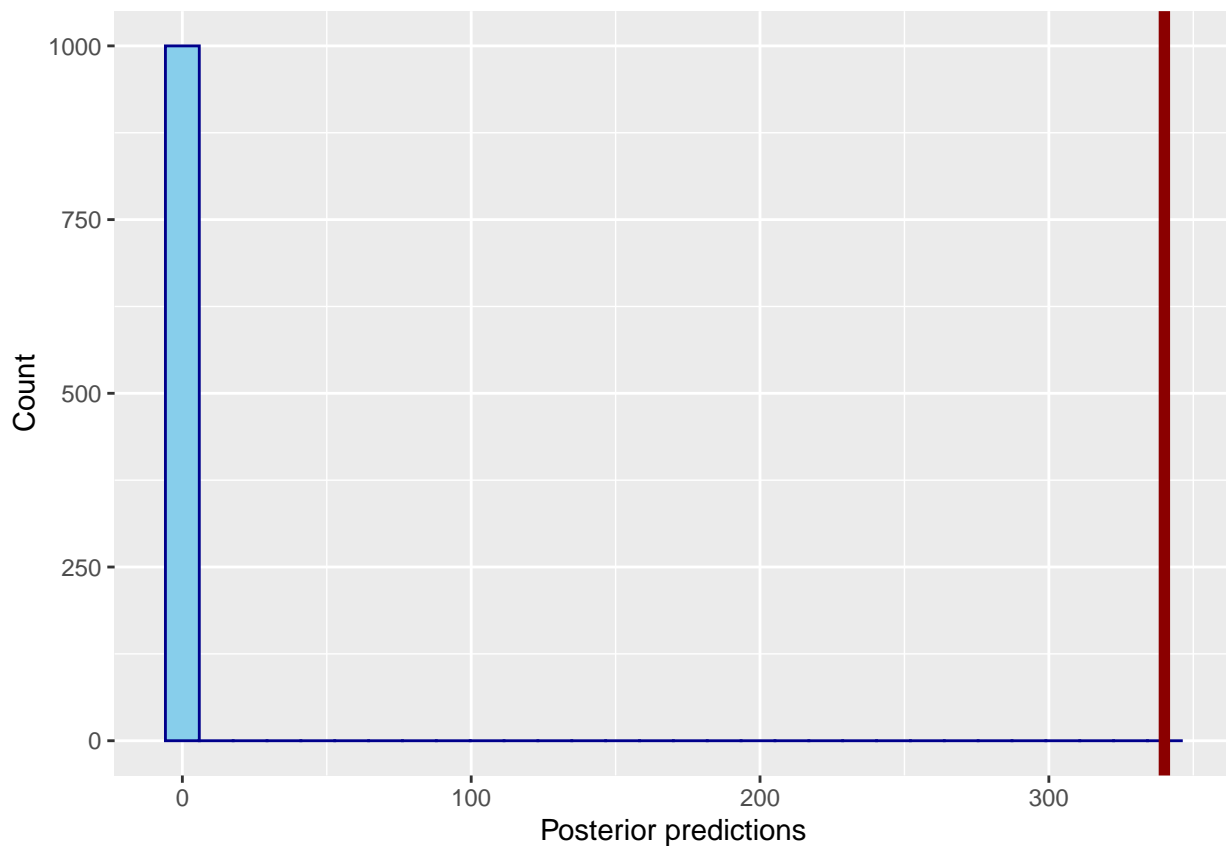
```
durham <- posterior_predict(m1,
  newdata = data.frame(pop = 316739,
    pct_rural = 6),
  draws = 1000)
head(durham)
```

```
##      1
```

```
## [1,] 1
## [2,] 1
## [3,] 0
## [4,] 1
## [5,] 1
## [6,] 1
```

```
ggplot(as.data.frame(durham), aes(x = durham)) +
  geom_histogram(color = "darkblue", fill = "skyblue") +
  labs(x = "Posterior predictions", y = "Count") +
  geom_vline(xintercept = 340, color = "darkred", lwd = 2)
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



Is there evidence of an association between rurality and being a high-crash county, after adjusting for population? Explain, and quantify any variability in your estimates. Be sure to evaluate model convergence.

asks the beta coefficient for pct_rural variability in beta_1